

# Security Audit Report


---

## Pokt Network Pokt Sparse Merkle Tree

Initial Report // February 20, 2024  
Final Report // April 4, 2024

### Team Members

Jehad Baeth // Senior Security Auditor  
Justin Regele // Senior Security Auditor  
Bashir Abu-Amr // Head of Delivery



## Table of Contents

1.0 Scope	— 2
↳ <a href="#">1.1 Technical Scope</a>	
↳ <a href="#">1.2 Documentation</a>	
↳ <a href="#">1.3 Bibliography</a>	
2.0 Executive Summary	— 3
↳ <a href="#">2.1 Schedule</a>	
↳ <a href="#">2.2 Overview</a>	
↳ <a href="#">2.3 Threat Model</a>	
↳ <a href="#">2.4 Security by Design</a>	
↳ <a href="#">2.5 Secure Implementation</a>	
↳ <a href="#">2.6 Use of Dependencies</a>	
↳ <a href="#">2.7 Tests</a>	
↳ <a href="#">2.8 Project Documentation</a>	
3.0 Key Findings Table	— 6
4.0 Findings	— 7
↳ <a href="#">4.1 Insufficiently Secure Hash Function Can Be Initialized</a>	
✓ Low    ✓ Fixed	
↳ <a href="#">4.2 Writeable KV Store Could Prove Non-Existent Nodes</a>	
✓ Low    ✓ Fixed	
↳ <a href="#">4.3 Improve Calling Convention of VerifyClosestProof and verifyProofWithUpdates</a>	
✓ None    ✓ Fixed	
↳ <a href="#">4.4 Refactor the Insertion of an Extension Node</a>	
✓ None    ✓ Fixed	
5.0 Appendix A	— 11
↳ <a href="#">5.1 Severity Rating Definitions</a>	
6.0 Appendix B	— 12
↳ <a href="#">6.1 Thesis Defense Disclaimer</a>	

# About Thesis Defense

Thesis Defense serves as the auditing services arm within Thesis, Inc., the venture studio behind tBTC, Fold, Tahoe, Etcher, and Mezo. Our team of security auditors have carried out hundreds of security audits for decentralized systems across a number of technologies including smart contracts, wallets and browser extensions, bridges, node implementations, cryptographic protocols, and dApps. We offer our services within a variety of ecosystems including Bitcoin, Ethereum + EVMs, Stacks, Cosmos / Cosmos SDK, NEAR and more.

Thesis Defense will employ the Thesis Defense Audit Approach and Audit Process to the in scope service. In the event that certain processes and methodologies are not applicable to the in scope services, we will indicate as such in individual audit or design review SOWs. In addition, Thesis Defense provides clear guidance on successful Security Audit Preparation.

## Section 1.0

# Scope

## Technical Scope

- **Repository:** <https://github.com/pokt-network/smt>
- **Audit Commit:** 868237978c0b3c0e2added161b36eeb7a3dc93b0
- **Verification Commit:** 3981639bd08cf52a7668c3681cbe2243d957e4ee

Any third-party or dependency library code is considered out of scope, unless explicitly specified as in scope above.

## Documentation

- Technical documentation: <https://github.com/pokt-network/smt/tree/main/docs>

## Bibliography

- M. Al-Bassam, A. Sonnino, and V. Buterin, "Fraud proofs: Maximising light client security and scaling blockchains with dishonest majorities," arXiv preprint arXiv:1809.09044, vol. 160, 2018.
- M. Al-Bassam, "Lazyledger: A distributed data availability ledger with client-side smart contracts," arXiv preprint arXiv:1905.09274, 2019.
- R. Dahlberg, T. Pulls, and R. Peeters, "Efficient Sparse Merkle Trees: Caching strategies and secure (non-) membership proofs," in Secure IT Systems: 21st Nordic Conference, NordSec 2016, Oulu, Finland, November 2-4, 2016. Proceedings 21, 2016, pp. 199–215.
- Z. Gao, Y. Hu, and Q. Wu, "Jellyfish Merkle tree," 2021.
- Z. Amsden et al., "The Libra Blockchain - White Paper," Whitepaper, 2019.
- D. Olshansky and R. R. Colmeiro, "Relay Mining: Verifiable Multi-Tenant Distributed Rate Limiting," arXiv preprint arXiv:2305.10672, 2023.
- J. Kalidhindi, A. Kazorian, A. Khera, and C. Pari, "Angela: A sparse, distributed, and highly concurrent merkle tree," UC Berkeley, Berkeley, 2018.



# Executive Summary

---

## Schedule

This security audit was conducted from February 7, 2024 to February 20, 2024 by 2 security auditors for a total of 4 person-weeks.

## Overview

Thesis Defense conducted a manual code review of POKT Network's Sparse Merkle Trie (Trie) implementation, an authenticated key-value data structure.

The POKT Network team has adapted a Sparse Merkle Tree for POKT Network's use case. POKT Network nodes provide applications with RPC relay services. Upon claiming rewards, the node must prove that it provided the work of the service by generating a valid inclusion "closest proof" of a randomly selected branch within the agreed upon root trie.

The SMT algorithm is well tested and mature. The scope of this audit included changes that have been made to improve the efficiency of the algorithm.

## Threat Model

For this audit, our team considered a threat model whereby a malicious node is able to generate a valid proof of inclusion to earn illegitimate rewards. We also considered an honest node whose valid proof of inclusion is rejected.

## Security by Design

The POKT SMT is designed to efficiently handle sparse data where most of the tree's leaves are empty, and is particularly useful for key value stores, allowing for an efficient proof of non-inclusion. Each leaf corresponds to a key value pair, with that leaf's position predetermined by the hash of that leaf's key.

In a SHA256 trie, any update would take 256 steps, while generating a proof could use up to 256 steps.

## ProveClosest

We investigated the `ProveClosest` algorithm and found that it is a more efficient version of a non-inclusion proof. `ProveClosest` is a more efficient and specific proof that shows the non-inclusion of a key by demonstrating the emptiness of the closest leaf node, while a normal non-inclusion proof is a more general proof that demonstrates the non-inclusion of a key by showing that no path in the tree leads to that key. Closest proof makes use of the fact that in a SMT, the vast majority of nodes are empty or hold nil values.

## Secure Implementation

We found the code to be well organized and correctly implemented, in accordance with security best practices. We investigated the security of the implementation of the POKT SMT in several key areas.

## False Sum Proof Protection

We investigated whether a Sum Proof could be forged by manipulating the weight appended to the end of the leaf digest. We concluded that even if an attacker forged a leaf where a weight was hashed into the leaf digest, but a different weight was appended, the proof would fail in the `validateBasic` validation checks.



```

siblingHash := hashPreimage(spec, proof.SiblingData)
if eq := bytes.Equal(proof.SideNodes[0], siblingHash);
!eq {
    return fmt.Errorf("invalid sibling data hash: got %x but want %x", siblingHash, proof.SideNodes[0])
}

```

Here, `proof.SiblingData` is the sibling leaf node of the node being proven. An attacker can add arbitrary data to this value in the Proof, but this verification check will hash the entire data, and in the case of a Sum Proof, will append the weight at the end of the hash, even though the weight has been included in the preimage.

This essentially is enforcement for the weights inside the proof. This is effective because we can provide arbitrary `SiblingData` in the proof and we can have the hash of weight A in the `SiblingData` with weight B appended, but when `SiblingData` is hashed with weight B, then it will not equal the `SideNode` hash. If we leave weight as is in `SiblingData`, the hashes match, but not the weights, resulting in the following error:

```

invalid sibling data hash: got
ba36043b00bccd6730e474139979b2954573b5e3520581e5cb542
010163f81d30000000000000000a

```

but want

```

ba36043b00bccd6730e474139979b2954573b5e3520581e5cb542
010163f81d300000000000000005"

```

Note that the hashes themselves are identical, but the weights differ after being manipulated in proof construction.

## Path Representation

Representing the path in a Merkle tree can be done either as a series of bits indicating the node's location based on its position in the tree or as an index value. This method is straightforward and can be implemented efficiently, especially for binary Merkle trees. The binary representation of the index directly corresponds to the path from the root to the leaf node. Currently Paths in the SMT implementation are represented as a series of bits.

It also can be represented with `N(level, depth)` since depth of all leaf nodes is the same given a specific hash function a path can be represented with depth only. Using an index value to represent the path can be more intuitive and easier to understand, as it directly maps to the position of the leaf node in the tree. An index value is also more compact than a series of bits, as it does not require storing the entire path.

That being said, we think further exploration of the design decision trade-offs between efficiency, compactness, and ease of implementation could have positive impacts for future iterations.

## Parallelism in SMT

The current implementation of SMT does not account for parallelism. Literature concerning other implementations of an SMT utilized a variety of techniques such as parallel processing, atomic update operations, and batch processing for further optimizations. We did not find any indication that the code anticipates handling parallel processing.

We investigated the access patterns of existing byte arrays that hold different data used by the code. Optimally, access to such resources should be synchronized when there are multiple writers interacting with it at the same time. We have only been able to identify cases where a single writer interacts with those variables at a time. However, we believe that the SMT implementation can be encapsulated with goroutines for protection against unexpected concurrency edge cases.

Edit: During the verification phase of this security audit, the POKT Network team acknowledged this optimization as a feature that can be implemented during a future phase.



## Protection Against Second Preimage Attacks and Shortened Proof Attacks

A second preimage attack on a SMT occurs when an attacker can find two different inputs that produce the same hash output. In the context of Merkle trees, this could mean finding two different sets of data that, when hashed into the tree, result in the same root hash. Currently, the SMT uses node-type specific prefixes (0 for leaf nodes, 1 for inner nodes, and 2 for extension nodes) to protect against second preimage attacks, which is a valid strategy commonly utilized in SMT implementations.

Alternatively, shortened proof attacks are a type of security vulnerability that can occur in Merkle trees when a proof is not constructed correctly or when there is a lack of domain separation between leaf nodes and internal nodes. This type of attack is similar to a second preimage attack, where an attacker can create a valid proof for a data element that is not actually in the original key-value store. Using a different hash function for hashing leaves than for hashing internal nodes helps mitigate the risk of both attacks. We recommend the POKT Network team explore utilizing such a strategy to further enforce domain separation, which provides protection against both second preimage and shortened proof attacks.

## Use of Dependencies

We ran dependency analysis tools and did not identify any issues in the use of dependencies.

## Tests

There are tests implemented in the repository with benchmark and stress testing documented. These tests are helpful in understanding the intended functionality, and are generally sufficient, but narrow in scope.

Before deploying the new SMT implementation onto mainnet, we recommend running a shadow fork where developers can test the new implementation in a controlled, safe, and isolated environment. This will reduce the risk of disrupting the mainnet or causing unintended consequences, in addition to identifying any issues or bugs that may arise. Furthermore, running a shadow fork would allow conducting a more thorough performance evaluation under different conditions and loads. Lastly, shadow forks can help identify compatibility issues between the new SMT implementation and other system components.

## Project Documentation

There is comprehensive documentation available for the currently deployed POKT Network “Morse”. In addition, there is sufficient documentation detailing the research, reasoning, and rationale for the design choices made regarding POKT Network’s protocol “Shannon” upgrade. The code is well commented and the documentation available in the code repository is accurate and helpful.

In addition, we received updates to the technical documentation during the audit. Although we were not able to review the updated documentation thoroughly due to time constraints, the work on this documentation should continue and we recommend a thorough review of the updated documentation during the verification phase.



# Key Findings Table

Issues	Severity	Status
ISSUE #1 Insufficiently Secure Hash Function Can Be Initialized	✓ Low	✓ Fixed
ISSUE #2 Writeable KV Store Could Prove Non-Existent Nodes	✓ Low	✓ Fixed
ISSUE #3 Improve Calling Convention of <code>VerifyClosestProof</code> and <code>verifyProofWithUpdates</code>	⏏ None	✓ Fixed
ISSUE #4 Refactor the Insertion of an Extension Node	⏏ None	✓ Fixed

Severity definitions can be found in [Appendix A](#)



# Findings

We describe the security issues identified during the security audit, along with their potential impact. We also note areas for improvement and optimizations in accordance with best practices. This includes recommendations to mitigate or remediate the issues we identify, in addition to their status before and after the fix verification.

ISSUE#1

## Insufficiently Secure Hash Function Can Be Initialized

✓ Low

✓ Fixed

### Location

<https://github.com/pokt-network/smt/blob/868237978c0b3c0e2added161b36eeb7a3dc93b0/smt.go#L74C9-L74C18>

### Description

In Go, the hash package provides a standard interface for hash functions and checksum algorithms. The `hash.Hash` interface is implemented by all hash functions in the standard library. A hash function suitable for the Merkle tree implementation must satisfy some criteria such as:

- Collision Resistance: For a Merkle tree, a hash function that is collision-resistant is required. This means that it is computationally infeasible to find two different inputs that hash to the same output.
- Preimage Resistance: A hash function that is preimage-resistant means that it is difficult to find the original input given only the hash output. This property is crucial for the security of a Merkle tree, as it prevents an attacker from modifying the data without being detected.
- Efficiency: The hash function should be efficient, as it will be used to compute the hash of many nodes in the Merkle tree.

SMT poses no hardcoded restriction on the type of hash functions that can be used other than that it should implement go's hash interface. Nor the documentation provides guidance on which hash functions can be used.

### Impact

The implementation of an insufficiently collision or preimage resistant hash function could undermine the security assumptions of the SMT.

### Recommendation

We recommend implementing a code level hard restriction on what hash functions are allowed to initialize the SMT based on collision, preimage attack, and time and space efficiency.

### Verification Status

The Pokt SMT documentation has been updated to explicitly explain the criteria of a hash function selection.



#### ISSUE#2

## Writeable KV Store Could Prove Non-Existent Nodes

✓ Low

✓ Fixed

### Location

[kvstore/interfaces.go#L12](#)

### Description

A Malicious Prover with write access to the Key Value (KV) store of a Verifier could write arbitrary data over a Leaf Node in order to convince the SMT that the node is an Internal Node, with the intention of writing Proofs for nodes that do not exist. This attack would only work if the Verifier updated the sibling leaf node of the corrupted data, which would update the Merkle Root to reflect the presence of the malicious internal node, allowing the attacker to prove the existence of child nodes of the internal node, without them existing in the SMT.

### Impact

The integrity of the SMT's Proving mechanism can be compromised.

### Recommendation

We recommend ensuring that the KV store is not writeable externally. The POKT Network team confirmed that the KV store is not replicated.

### Verification Status

The Pokt SMT documentation explicitly warns users to ensure usage of an externally writable key-value store in production environments.

#### ISSUE#3

## Improve Calling Convention of `VerifyClosestProof` and `verifyProofWithUpdates`

None

✓ Fixed

### Location

[proofs.go#L320, L323](#)

### Description

The way the `verifyProofWithUpdates` looks up Paths by hashing the Key creates an awkward programming interface with the `VerifyClosestProof` algorithm, because a `ClosestProof` provides a Path as a Key. If a new Spec of a `nilPathHasher` is not provided during Proof Verification, double hashing of the Path will occur and the Proof will fail.

### Impact

None – no security impact

### Recommendation

We recommend revising the calling conventions around `VerifyClosestProof` and `verifyProofWithUpdates` such that the API is not implemented as a workaround.





## Verification Status

The updated code encapsulates the creation of the `nilPathHasher` so that the calling convention remains standard across the application.

ISSUE#4

## Refactor the Insertion of an Extension Node

None

Fixed

## Location

[smt.go#L179-L189](#)

## Description

The insertion of Extension Nodes uses the same pointer to assign values to Extension Nodes as well as its children. We found this coding pattern confusing.

## Impact

None – no security impact.

## Recommendation

We recommend updating the Extension Node's child with an explicit variable as in the example below, leaving the pointer variable to only dereference the Node at the current depth.



```

var extension_node_child = nil
if getPathBit(path, prefixlen) == left {
    extension_node_child = & innerNode {
        leftChild: newLeaf,
        rightChild: leaf
    }
} else {
    extension_node_child = & innerNode {
        leftChild: leaf,
        rightChild: newLeaf
    }
}
ext: = extensionNode {
    child: extension_node_child,
    path: path,
    pathBounds: [2] byte {
        byte(depth), byte(prefixlen)
    }
} * last = & ext

```

## Verification Status






The POKT Network team has altered the mentioned coding pattern and currently employs separate pointers, specifically for the Child nodes.



# Appendix A

## Severity Rating Definitions

At Thesis Defense, we utilize the [Immunefi Vulnerability Severity Classification System - v2.3](#).

Severity	Definition
 Critical	<ul style="list-style-type: none"> <li>• Manipulation of governance voting result deviating from voted outcome and resulting in a direct change from intended effect of original results</li> <li>• Direct theft of any user funds, whether at-rest or in-motion, other than unclaimed yield</li> <li>• Direct theft of any user NFTs, whether at-rest or in-motion, other than unclaimed royalties</li> <li>• Permanent freezing of funds</li> <li>• Permanent freezing of NFTs</li> <li>• Unauthorized minting of NFTs</li> <li>• Predictable or manipulable RNG that results in abuse of the principal or NFT</li> <li>• Unintended alteration of what the NFT represents (e.g. token URI, payload, artistic content)</li> <li>• Protocol insolvency</li> </ul>
 High	<ul style="list-style-type: none"> <li>• Theft of unclaimed yield</li> <li>• Theft of unclaimed royalties</li> <li>• Permanent freezing of unclaimed yield</li> <li>• Permanent freezing of unclaimed royalties</li> <li>• Temporary freezing of funds</li> <li>• Temporary freezing NFTs</li> </ul>
 Medium	<ul style="list-style-type: none"> <li>• Smart contract unable to operate due to lack of token funds</li> <li>• Enabling/disabling notifications</li> <li>• Griefing (e.g. no profit motive for an attacker, but damage to the users or the protocol)</li> <li>• Theft of gas</li> <li>• Unbounded gas consumption</li> </ul>
 Low	<ul style="list-style-type: none"> <li>• Contract fails to deliver promised returns, but doesn't lose value</li> </ul>
 None	<ul style="list-style-type: none"> <li>• We make note of issues of no severity that reflect best practice recommendations or opportunities for optimization, including, but not limited to, gas optimization, the divergence from standard coding practices, code readability issues, the incorrect use of dependencies, insufficient test coverage, or the absence of documentation or code comments.</li> </ul>



# Appendix B

---

## Thesis Defense Disclaimer

Thesis Defense conducts its security audits and other services provided based on agreed-upon and specific scopes of work (SOWs) with our Customers. The analysis provided in our reports is based solely on the information available and the state of the systems at the time of review. While Thesis Defense strives to provide thorough and accurate analysis, our reports do not constitute a guarantee of the project's security and should not be interpreted as assurances of error-free or risk-free project operations. It is imperative to acknowledge that all technological evaluations are inherently subject to risks and uncertainties due to the emergent nature of cryptographic technologies.

Our reports are not intended to be utilized as financial, investment, legal, tax, or regulatory advice, nor should they be perceived as an endorsement of any particular technology or project. No third party should rely on these reports for the purpose of making investment decisions or consider them as a guarantee of project security.

Links to external websites and references to third-party information within our reports are provided solely for the user's convenience. Thesis Defense does not control, endorse, or assume responsibility for the content or privacy practices of any linked external sites. Users should exercise caution and independently verify any information obtained from third-party sources.

The contents of our reports, including methodologies, data analysis, and conclusions, are the proprietary intellectual property of Thesis Defense and are provided exclusively for the specified use of our Customers. Unauthorized disclosure, reproduction, or distribution of this material is strictly prohibited unless explicitly authorized by Thesis Defense. Thesis Defense does not assume any obligation to update the information contained within our reports post-publication, nor do we owe a duty to any third party by virtue of making these analyses available.

